

FROM SOURCE EVIDENCE TO PUBLIC ARCHIVE

How a Record becomes searchable, classified, and connected

A landscape briefing deck for explaining intake, classification, tagging, entity extraction, relationship mapping, and the static frontend delivery path.

Canonical Record

IDENTITY	id, title, URL, author, date
PLACEMENT	Record type, era, publication, categories
MEANING	summary, pull quote, concepts, notes
CONNECTIONS	Entities, typed relationships, relatedIds
DELIVERY	CSV to JSON to Loader to Entity Index

INTAKE

NORMALIZE

CLASSIFY

MAP

PUBLISH

CURATOR CONTROL

CSV row, verification, canonical URL, source notes, and final judgment.

AI ASSISTANCE

Drafted summaries, concept suggestions, entity extraction, and relationship proposals.

QUALITY GATES

Schema checks, deduplication, referential integrity, and source recovery tracking.

PUBLIC SURFACES

Search, Facets, Record modal, Explorer, Entity Browser, and Analytics.

1,030

MAIN
RECORDS

29,700

SOCIAL POSTS

5,036

ENTITIES

4,666

RELATIONSHIPS

Six stages for every Record

Each source moves through a fixed chain: capture the evidence, normalize the Record, classify it for browsing, add interpretive fields, map Entities and relationships, then publish static data for the frontend.

1 Intake

Capture URL, clipping, thread, transcript, Tumblr post, or source file. Record provenance and recovery status.

INPUT

Source evidence

CHECK

Recoverability and verification

OUTPUT

Candidate Record row

source URL

raw text

verified status

2 Normalize

Convert messy source material into the canonical Record shape used by the Archive data files.

INPUT

Extracted metadata

CHECK

ID and date format

OUTPUT

Clean Record fields

id

date

publication

3 Classify

Assign the browsing fields that place the Record by format, time period, publication, and category.

INPUT

Record text and source

CHECK

Taxonomy spelling

OUTPUT

Facet-ready fields

Record type

era

categories

4 Tag

Add the readable editorial layer: summary, pull quote, concepts, notes, and legacy tags.

INPUT

Raw text and context

CHECK

Curator review

OUTPUT

Readable detail view

summary

pull quote

concepts

5 Map

Extract named Entities, deduplicate them, and attach typed relationships to the Record.

INPUT

Record text

CHECK

Entity registry match

OUTPUT

Graph connections

people

organizations

relationships

6 Publish

Generate static JSON. The frontend Loader builds Facets, search, modals, Explorer, and the Entity Index.

INPUT

Canonical CSVs

CHECK

Export and data tests

OUTPUT

Static Archive data

archive-core

archive-details

archive-entities

Evidence capture and canonical fields

Curator lane

CSV source of truth

Pipeline lane

Extraction and repair

1 Choose the source

PressThink essay, external article, Tumblr archive, social thread, audio/video transcript, or clipping.

source scope

2 Create the Record ID

Use the next canonical ID, such as RECORD-00902 for a new main Record or BSKY-03122 for a new Bluesky post.

ID policy sequence

3 Enter baseline metadata

Title, URL, author, publication date, original publication, content type, format, and platform.

metadata CSV row

4 Record provenance

Verification flag, source notes, canonical URL choice, recovery status, and any constraints around the original.

verified notes

5 Save the canonical row

The CSVs under `data/` remain the version-controlled input for the static Archive export.

canonical versioned

1 Route by source type

Article, audio, video, Twitter/X, Bluesky, Tumblr, and clipping processors use source-specific strategies.

dispatcher processor

2 Extract usable text

Gemini URL Context is the first path for web articles. Fallbacks include requests, BeautifulSoup, Playwright, transcripts, and OCR.

raw_text fallbacks

3 Clean and structure

Remove boilerplate, preserve paragraphs, capture `raw_text`, and keep source-specific fields available for later analysis.

clean text fields

4 Backfill gaps

Date backfills, publisher normalization, raw text recovery, pull quote extraction, and content quality scoring improve older Records.

repair quality

5 Validate the row

Checks catch bad IDs, missing dates, duplicate URLs, broken references, empty raw text, and schema drift.

tests schema

The shelves, facets, and historical placement

Record catalog fields

These fields drive sorting, filters, cards, deep links, and editorial review.

RECORD TYPE

RECORD, TUMBLR,
CLIP, THREAD,
Dissertation.

CONTENT TYPE

Article, essay, interview,
lecture, social media post,
and related forms.

FORMAT

Text, audio, video,
clipping, transcript, or
thread structure.

PUBLICATION

PressThink, external
outlets, platforms, and
source domains.

ERA

Historical period used for
browsing and narrative
context.

CATEGORIES

One or more thematic
labels used as Facets.

VERIFIED

Curator confirmation of
source quality and
Archive readiness.

CANONICAL URL

Modern URL, archive
URL, or recovered source
path.

RELATED IDS

Record and Entity
references used by detail
views.

A Press & Media Criticism

Coverage of journalism performance, press behavior, media institutions, and public trust.

B Journalism Theory & Practice

Core professional concepts, methods, norms, and debates about how journalism works.

C Politics & Democracy

Records connecting journalism to elections, publics, civic life, legitimacy, and democratic practice.

D Technology & Digital Media

Blogging, platforms, networks, audience change, social media, and digital publishing.

E Audience & Public Engagement

Citizens, participation, listening, public knowledge, and audience relationships.

F Journalism Education

Teaching, curriculum, student-facing work, and Rosen's NYU role.

Summary, pull quote, concepts, and notes

Summary

A short explanation of the Record's argument, context, or role in the Archive. This supports scan-friendly cards, modals, and editorial QA.

SOURCE

raw_text, title, date, publication, source notes.

PROCESSING

AI draft plus curator tightening for specificity and tone.

FRONTEND USE

Record cards, search context, modal lead text, review scans.

Use: Tell Jay what this Record contributes in one tight paragraph, with enough specificity to distinguish it from nearby Records.

Pull quote

A representative excerpt that carries the Record's voice. It helps readers recognize the source and gives the modal a strong entry point.

SOURCE

verified source text, transcript, OCR, or recovered article body.

PROCESSING

Exact excerpt selection, quote hygiene, length control.

FRONTEND USE

Detail modal, featured Records, editorial review, share copy.

Use: Preserve exact wording when quoting source text. Keep the quote attached to source provenance and review status.

Concept strings

Record-level idea labels: View from Nowhere, Citizens' Agenda, Mindcasting, Audience Atomization Overcome, The Production of Innocence, and related Rosen vocabulary.

SOURCE

Record argument, recurring Rosen vocabulary, dissertation-era themes.

PROCESSING


Schema-constrained concept assignment and normalization.

FRONTEND USE

Modal context, filtering candidates, future concept pages.

Use: Keep these labels distinct from structured Concept Entities. The labels describe the Record; Entities power the graph.

The knowledge graph layer



CSV to JSON to Archive experience


1 Canonical CSVs
archive_records-public.csv, social_posts.csv,
extracted_entities.csv, extracted_relationships.csv.

2 Export script
node data/export-archive-data.js builds processed
Records, Entities, Facets, and autocomplete terms.

3 Split JSON
archive-core loads first. archive-details and archive-
entities load on demand. archive-data remains the
fallback.

4 Loader
The frontend fetches static JSON over HTTP and
handles browser caching.

5 Public views
Archive browser, Record modal, Entity Browser,
Explorer, Analytics, and Dissertation tools.



Talk track for Jay

1

Start with the Record

The Record is the stable archival unit. Every feature starts from that row: evidence, metadata, classification, editorial fields, and relationships.

2

Separate fields by job

Identity fields cite the source. Classification fields power browsing. Tagging fields explain meaning. Entity fields create cross-Archive connections.

3

Use AI as processing support

AI drafts categories, summaries, concepts, and extractions. Schema validation, deduplication, and curator review keep the Archive coherent.

4

Show the frontend payoff

Static JSON gives the browser enough structure to support search, Facets, Record modals, the Entity Browser, and network exploration.

Each Record carries evidence, interpretation, and connections.

The Archive workflow turns source material into a durable public reference point, then links that reference point to the larger map of Rosen's people, institutions, ideas, works, and eras.